

Advanced Machine Learning

Cristian Perez Jensen

January 17, 2025

Note that these are not the official lecture notes of the course, but only notes written by a student of the course. As such, there might be mistakes. The source code can be found at github.com/cristianpjensen/eth-cs-notes. If you find a mistake, please create an issue or open a pull request.

Contents

1	<i>Paradigms of data science</i>	1
2	<i>Anomaly detection</i>	2
3	<i>Density estimation</i>	6
4	<i>Regression</i>	9
4.1	<i>Linear regression</i>	9
4.2	<i>Polynomial regression</i>	10
4.3	<i>Regularization</i>	11
5	<i>Causality</i>	12
5.1	<i>Counterfactual invariance</i>	12
6	<i>Gaussian processes</i>	14
6.1	<i>Kernels</i>	15
6.2	<i>Prediction</i>	16
7	<i>Uncertainty quantification</i>	17
7.1	<i>Statistical model validation</i>	17
7.2	<i>Uncertainty in linear models</i>	19
7.3	<i>Statistical testing</i>	19
7.4	<i>Bayesian neural networks</i>	20
7.5	<i>Information-based transductive learning</i>	22
8	<i>Convex optimization</i>	25
8.1	<i>Support vector machine</i>	25
9	<i>Ensembles</i>	26
9.1	<i>Bagging</i>	26
9.2	<i>Random forests</i>	26
9.3	<i>AdaBoost</i>	26
10	<i>Stable diffusion</i>	27
10.1	<i>Diffusion models</i>	27
10.2	<i>U-net</i>	28
10.3	<i>Latent diffusion models</i>	28
10.4	<i>Text embeddings</i>	29
10.5	<i>Cross-attention</i>	29

List of symbols

\doteq	Equality by definition
\approx	Approximate equality
\propto	Proportional to
\mathbb{N}	Set of natural numbers
\mathbb{R}	Set of real numbers
$i : j$	Set of natural numbers between i and j . <i>I.e.</i> , $\{i, i+1, \dots, j\}$
$f : A \rightarrow B$	Function f that maps elements of set A to elements of set B
$\mathbb{1}\{\text{predicate}\}$	Indicator function (1 if predicate is true, otherwise 0)
$\boldsymbol{v} \in \mathbb{R}^n$	n -dimensional vector
$\boldsymbol{M} \in \mathbb{R}^{m \times n}$	$m \times n$ matrix
\boldsymbol{M}^\top	Transpose of matrix \boldsymbol{M}
\boldsymbol{M}^{-1}	Inverse of matrix \boldsymbol{M}
$\det(\boldsymbol{M})$	Determinant of \boldsymbol{M}
$\frac{\mathrm{d}}{\mathrm{d}x}f(x)$	Ordinary derivative of $f(x)$ w.r.t. x at point $x \in \mathbb{R}$
$\frac{\partial}{\partial \boldsymbol{x}}f(\boldsymbol{x})$	Partial derivative of $f(\boldsymbol{x})$ w.r.t. \boldsymbol{x} at point $\boldsymbol{x} \in \mathbb{R}^n$
$\nabla_{\boldsymbol{x}}f(\boldsymbol{x}) \in \mathbb{R}^n$	Gradient of $f : \mathbb{R}^n \rightarrow \mathbb{R}$ at point $\boldsymbol{x} \in \mathbb{R}^n$
$\nabla_{\boldsymbol{x}}^2f(\boldsymbol{x}) \in \mathbb{R}^{n \times n}$	Hessian of $f : \mathbb{R}^n \rightarrow \mathbb{R}$ at point $\boldsymbol{x} \in \mathbb{R}^n$

1 Paradigms of data science

Let $\{x_1, \dots, x_n\}$ be i.i.d. samples, generated by an unknown distribution P . Assume that this distribution is in a distribution family,

$$\mathcal{H} = \{p(\cdot \mid \theta) \mid \theta \in \Theta\}.$$

The goal is to learn the parameters θ that fit the data $\{x_1, \dots, x_n\}$ best.

Frequentism. In frequentism, the maximum likelihood estimator (MLE) parameters maximize the following,

$$\begin{aligned} \theta^* &\in \operatorname{argmax}_{\theta \in \Theta} \log p(\{x_1, \dots, x_n\} \mid \theta) \\ &= \operatorname{argmax}_{\theta \in \Theta} \sum_{i=1}^n \log p(x_i \mid \theta). \end{aligned}$$

Bayesianism. Bayesianism assumes that there is a prior over distributions. The maximum a posteriori (MAP) parameters maximize the following,

$$\begin{aligned} \theta^* &\in \operatorname{argmax}_{\theta \in \Theta} \log p(\theta \mid \mathbf{X}) \\ &= \operatorname{argmax}_{\theta \in \Theta} \log p(\{x_1, \dots, x_n\} \mid \theta) \cdot p(\theta) && \text{Bayes' rule.} \\ &= \operatorname{argmax}_{\theta \in \Theta} \log p(\theta) + \sum_{i=1}^n \log p(x_i \mid \theta). \end{aligned}$$

In practice, the prior acts as a regularization term.

Statistical learning. Now, assume that we have labeled samples $\{(x_1, y_1), \dots, (x_n, y_n)\} \subseteq \mathcal{X} \times \mathcal{Y}$, where y is the target variable. Let $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ be a loss function. For a predictor function $f : \mathcal{X} \rightarrow \mathcal{Y}$, we define its risk as the expected loss,

$$\mathcal{R}(f) \doteq \mathbb{E}_{X,Y}[\ell(y, f(x))].$$

In statistical learning, we want to find a function that minimizes the risk. However, since the distribution over X, Y is unknown, we cannot compute $\mathcal{R}(f)$ directly. Instead, we use the empirical risk as a surrogate,

$$\hat{\mathcal{R}}(f) \doteq \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)).$$

The goal is to obtain the empirical risk minimizer,

$$f^* \in \operatorname{argmin}_{f \in \mathcal{F}} \hat{\mathcal{R}}(f),$$

where \mathcal{F} is a family of functions that we assume f belongs to.

2 Anomaly detection

In anomaly detection, we are given a sample of objects $\mathcal{X} \subseteq \mathbb{R}^d$ with a normal class $\mathcal{N} \in \mathcal{X}$ —the data points that we wish to keep. We wish to construct a function $\phi : \mathcal{X} \rightarrow \{0, 1\}$, such that

$$\phi(x) = 1 \iff x \notin \mathcal{N}.$$

Formally, an anomaly is an unlikely event. Hence, the strategy is to fit a model of a parametric family of distributions to the data \mathcal{X} ,

$$\mathcal{H} = \{p(\cdot | \theta) | \theta \in \Theta\}.$$

Then, we define the anomaly score of x as a low probability $p(x | \theta^*)$ according to the optimal model in this hypothesis class.

Anomaly detection in a high-dimensional space is hard, because the normal class can be very complex. The idea is to project \mathcal{X} down to a lower dimensionality and perform anomaly detection there—hopefully the projected version of the normal class $\Pi(\mathcal{N})$ is less complex. In order to find the optimal linear projection, we will use principal component analysis (PCA).

Furthermore, it has been observed that linear projections of high-dimensional distributions onto low-dimensional spaces resemble Gaussian distributions. Hence, after performing PCA, we will fit a Gaussian mixture model (GMM) to the projected data.

Principal component analysis. The goal of PCA is to linearly project \mathbb{R}^d to \mathbb{R}^{d^-} such that the maximum amount of variance of the data is preserved.¹ Consider the base case $d^- = 1$. Let $\mathbf{u} \in \mathbb{R}^d$ with $\|\mathbf{u}\| = 1$, we project onto \mathbf{u} by inner product,

$$x \mapsto \mathbf{u}^\top x.$$

The sample mean of the reduced dataset is computed by

$$\frac{1}{n} \sum_{i=1}^n \mathbf{u}^\top x_i = \mathbf{u}^\top \bar{x},$$

where \bar{x} is the sample mean of the original dataset. Further, the sample variance of the reduced dataset is

$$\begin{aligned} \frac{1}{n} \sum_{i=1}^n (\mathbf{u}^\top x_i - \mathbf{u}^\top \bar{x})^2 &= \frac{1}{n} \sum_{i=1}^n \mathbf{u}^\top (x_i - \bar{x})(x_i - \bar{x})^\top \mathbf{u} \\ &= \mathbf{u}^\top \left(\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(x_i - \bar{x})^\top \right) \mathbf{u} \\ &= \mathbf{u}^\top \Sigma \mathbf{u}, \end{aligned}$$

where Σ is the covariance matrix of the dataset. Since we want the projection that preserves the maximum variance, we have the following objective,

$$\mathbf{u}^* \in \operatorname{argmax}_{\|\mathbf{u}\|=1} \mathbf{u}^\top \Sigma \mathbf{u}.$$

¹ Components with larger variance are more informative.

The Lagrangian of this problem is

$$\mathcal{L}(\mathbf{u}; \lambda) = \mathbf{u}^\top \mathbf{\Sigma} \mathbf{u} + \lambda (1 - \|\mathbf{u}\|^2)$$

with gradient

$$\frac{\partial \mathcal{L}(\mathbf{u}; \lambda)}{\partial \mathbf{u}} = 2\mathbf{\Sigma} \mathbf{u} - 2\lambda \mathbf{u} \stackrel{!}{=} 0.$$

So, \mathbf{u} must satisfy $\mathbf{\Sigma} \mathbf{u} = \lambda \mathbf{u}$ — \mathbf{u} is an eigenvector of $\mathbf{\Sigma}$. It is easy to see that this must be the principal eigenvector by rewriting the objective,

$$\begin{aligned} \mathbf{u}^* &\in \operatorname{argmax}_{\|\mathbf{u}\|=1} \mathbf{u}^\top \mathbf{\Sigma} \mathbf{u} \\ &= \operatorname{argmax}_{\substack{\|\mathbf{u}\|=1 \\ (\mathbf{u}, \lambda) \in \operatorname{eig}(\mathbf{\Sigma})}} \lambda \|\mathbf{u}\|^2 \\ &= \operatorname{argmax}_{\substack{\mathbf{u} \in \mathbb{R}^d \\ (\mathbf{u}, \lambda) \in \operatorname{eig}(\mathbf{\Sigma})}} \lambda \\ &= \mathbf{u}_1. \end{aligned}$$

For $d^- > 1$, the remaining principal components can be computed with a similar idea. Iteratively, we factor out the previous principal components and do as above on the transformed dataset. For example, to get the second principal component, we first factor out the first principal component,

$$\mathcal{X}_1 \doteq \{\mathbf{x} - \operatorname{proj}_{\mathbf{u}_1}(\mathbf{x}) \mid \mathbf{x} \in \mathcal{X}\} = \{\mathbf{x} - \mathbf{u}_1^\top \mathbf{x} \cdot \mathbf{u}_1 \mid \mathbf{x} \in \mathcal{X}\}.$$

Then, we do the same as above.

Gaussian mixture model. The probability density function (PDF) of a Gaussian mixture model with k components is formalized as a convex combination of Gaussians,

$$p(\mathbf{x}; \boldsymbol{\theta}) = \sum_{j=1}^k \pi_j \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_j, \mathbf{\Sigma}_j).$$

The parameters of this model are

$$\boldsymbol{\theta} = \{\pi_j, \boldsymbol{\mu}_j, \mathbf{\Sigma}_j \mid j \in [k]\},$$

where $\sum_{j=1}^k \pi_j = 1$ and $\{\mathbf{\Sigma}_j \mid j \in [k]\}$ are positive definite. We fit the parameters of this model by maximizing the log-likelihood,

$$\begin{aligned} \boldsymbol{\theta}^* &\in \operatorname{argmax}_{\boldsymbol{\theta} \in \Theta} \log p(\{\mathbf{x}_1, \dots, \mathbf{x}_n\}; \boldsymbol{\theta}) \\ &= \operatorname{argmax}_{\boldsymbol{\theta} \in \Theta} \sum_{i=1}^n \log p(\mathbf{x}_i; \boldsymbol{\theta}) \\ &= \operatorname{argmax}_{\boldsymbol{\theta} \in \Theta} \sum_{i=1}^n \log \sum_{j=1}^k \pi_j \mathcal{N}(\mathbf{x}_i; \boldsymbol{\mu}_j, \mathbf{\Sigma}_j). \end{aligned}$$

```

1: Initialize  $\theta_0$ 
2: for  $t \in [T]$  do
3:    $q^* \in \operatorname{argmin}_q E(q, \theta_{t-1})$ 
4:    $\theta_t \in \operatorname{argmax}_\theta M(q^*, \theta)$ 
5: end for
6: return  $\theta_T$ 

```

Note that the above is intractable, so we would like to decompose it into tractable terms that can be computed. Let's assume that we know from which latent variable each data point was generated, then we can compute the MLE of the extended dataset $\{(x_i, z_i) \mid i \in [n]\}$ as

$$\begin{aligned}
\log p(\mathbf{X}, \mathbf{z}; \theta) &= \sum_{i=1}^n \log p(x_i, z_i) \\
&= \sum_{i=1}^n \log(\pi_{z_i} \mathcal{N}(x_i; \mu_{z_i}, \Sigma_{z_i})) \\
&= \sum_{i=1}^n \log \pi_{z_i} + \sum_{i=1}^n \log \mathcal{N}(x_i; \mu_{z_i}, \Sigma_{z_i}),
\end{aligned}$$

which is tractable to maximize. Let q be a distribution over $[k]$, then we can rewrite the log-likelihood into tractable terms,

$$\begin{aligned}
\log p(\mathbf{X}; \theta) &= \mathbb{E}_{z \sim q} [\log p(\mathbf{X}; \theta)] \\
&= \mathbb{E}_{z \sim q} \left[\log \left(\frac{p(\mathbf{X}, z; \theta)}{p(z \mid \mathbf{X}; \theta)} \right) \right] \\
&= \mathbb{E}_{z \sim q} \left[\log \left(\frac{p(\mathbf{X}, z; \theta)}{p(z \mid \mathbf{X}; \theta)} \frac{q(z)}{q(z)} \right) \right] \\
&= \underbrace{\mathbb{E}_{z \sim q} \left[\log \frac{p(\mathbf{X}, z; \theta)}{q(z)} \right]}_{\doteq M(q, \theta)} + \underbrace{\mathbb{E}_{z \sim q} \left[\log \frac{q(z)}{p(z \mid \mathbf{X}; \theta)} \right]}_{\doteq E(q, \theta)}.
\end{aligned}$$

These terms have the following two properties,

$$\begin{aligned}
\log p(\mathbf{X}; \theta) &\geq M(q, \theta), \quad \forall q, \theta \\
\log p(\mathbf{X}; \theta) &= M(q^*, \theta), \quad q^* = p(\cdot \mid \mathbf{X}; \theta), \quad \forall \theta.
\end{aligned}$$

Hence, we can use $M(q^*, \theta)$ as an approximation of $\log p(\mathbf{X}; \theta)$ around θ .

Theorem 2.1 (EM algorithm convergence). Using the expectation-maximization algorithm, $\{\log p(x; \theta_t)\}_{t=0}^T$ is non-decreasing.

Proof. Given a data point x and current parameters θ , we have the following update,

$$\theta' \in \operatorname{argmax}_{\theta \in \Theta} M(q^*, \theta).$$

Algorithm 1. The expectation-maximization algorithm, where

$$\begin{aligned}
M(q, \theta) &\doteq \mathbb{E}_{z \sim q} \left[\log \frac{p(\mathbf{X}, z; \theta)}{q(z)} \right] \\
E(q, \theta) &\doteq \mathbb{E}_{z \sim q} \left[\log \frac{q(z)}{p(z \mid \mathbf{X}; \theta)} \right].
\end{aligned}$$

Hence, we have

$$\log p(\mathbf{x}) = M(q^*, \boldsymbol{\theta}) \leq M(q^*, \boldsymbol{\theta}') \leq \log p(\mathbf{x}; \boldsymbol{\theta}').$$

Thus, $\{\log p(\mathbf{x}; \boldsymbol{\theta}_t)\}_{t=0}^T$ is non-decreasing. ■

Summary. In conclusion, given a set of data points \mathcal{X} with normal points $\mathcal{N} \subseteq \mathcal{X}$, we train an anomaly detector as follows,

1. Fit a projector $\pi : \mathbb{R}^d \rightarrow \mathbb{R}^{d^-}$ using PCA;
2. Fit a probability density function $p(\cdot \mid \boldsymbol{\theta})$ with k components to $\{\pi(\mathbf{x}) \mid \mathbf{x} \in \mathcal{X}\}$ using the EM algorithm;
3. For a point $\mathbf{x} \in \mathcal{X}$, its “anomaly score” is computed by $-\log p(\pi(\mathbf{x}); \boldsymbol{\theta})$.

3 Density estimation

In this section, we will consider parametric models,

$$\{p(\mathbf{x}; \boldsymbol{\theta}) \mid \boldsymbol{\theta} \in \Theta\}.$$

The problem we concern ourselves with is finding the best parameter $\boldsymbol{\theta}$. The most common method of finding the best parameters is maximum likelihood estimation (MLE),

$$\begin{aligned} \hat{\boldsymbol{\theta}}_{\text{MLE}} &\in \operatorname{argmax}_{\boldsymbol{\theta} \in \Theta} \prod_{i=1}^n p(\mathbf{x}_i; \boldsymbol{\theta}) \\ &= \operatorname{argmin}_{\boldsymbol{\theta} \in \Theta} - \sum_{i=1}^n \log p(\mathbf{x}_i; \boldsymbol{\theta}). \end{aligned}$$

The following properties makes the MLE estimator attractive,

1. Consistency—in the limit of n , $\hat{\boldsymbol{\theta}}_{\text{MLE}}$ converges to the true parameter $\boldsymbol{\theta}^*$;
2. Equivariance—if $\hat{\boldsymbol{\theta}}$ is the MLE of $\boldsymbol{\theta}$, then $g(\hat{\boldsymbol{\theta}})$ is the MLE of $g(\boldsymbol{\theta})$ for any function g ;
3. Asymptotically normal—in the limit of n , $\hat{\boldsymbol{\theta}}_{\text{MLE}} - \boldsymbol{\theta} / \sqrt{n}$ converges to a random variable with distribution $\mathcal{N}(0, \mathcal{I}(\boldsymbol{\theta})^{-1})$, where \mathcal{I} is the Fisher information matrix;
4. Asymptotically efficient—in the limit of n , the MLE estimator has the smallest variance of all unbiased estimators.

We can understand the asymptotic efficiency (property 4) of estimators better by the Rao–Cramér bound, which provides a bound on the variance of the estimator. We will only consider the general scalar case, but it generalizes to the multivariate case.

Theorem 3.1 (Rao–Cramér bound (scalar case)). For any (scalar) unbiased estimator $\hat{\theta} : \mathcal{Y}^n \rightarrow \mathbb{R}$, given n data points, of $\theta \in \mathbb{R}$, we have the following bound on its variance,

$$\operatorname{Var}[\hat{\theta}(\mathbf{y})] \geq \frac{\left(\frac{\partial}{\partial \theta} b_{\hat{\theta}} + 1\right)^2}{\mathcal{I}_n(\theta)} + b_{\hat{\theta}}^2,$$

where $\mathcal{I}_n(\theta)$ is the Fisher information,

$$\mathcal{I}_n(\theta) \doteq \mathbb{E}_{\mathbf{y}|\theta} \left[\left(\frac{\partial}{\partial \theta} \log p(\mathbf{y} \mid \theta) \right)^2 \right] \stackrel{\text{iid}}{=} n \cdot \mathcal{I}(\theta),$$

and $b_{\hat{\theta}}$ is the bias of $\hat{\theta}$,

$$b_{\hat{\theta}} \doteq \mathbb{E}_{\mathbf{y}|\theta} [\hat{\theta}(\mathbf{y})] - \theta.$$

Proof. Let the “score” be defined as follows,

$$\Lambda(\mathbf{y}, \theta) \doteq \frac{\partial}{\partial \theta} \log p(\mathbf{y} | \theta) = \frac{1}{p(\mathbf{y} | \theta)} \frac{\partial}{\partial \theta} p(\mathbf{y} | \theta).$$

Then we have $\mathcal{I}_n(\theta) = \mathbb{E}_{\mathbf{y}|\theta} [\Lambda(\mathbf{y}, \theta)^2]$.

The expected score is equal to zero,

$$\begin{aligned} \mathbb{E}_{\mathbf{y}|\theta} [\Lambda(\mathbf{y}, \theta)] &= \int p(\mathbf{y} | \theta) \Lambda(\mathbf{y}, \theta) d\mathbf{y} \\ &= \int \frac{\partial}{\partial \theta} p(\mathbf{y} | \theta) d\mathbf{y} \\ &= \frac{\partial}{\partial \theta} \int p(\mathbf{y} | \theta) d\mathbf{y} \\ &= \frac{\partial}{\partial \theta} 1 \\ &= 0. \end{aligned}$$

(Hence, the Fisher information is equivalent to the variance of the score.)

Furthermore, the cross-correlation between $\Lambda(\mathbf{y}, \theta)$ and $\hat{\theta}(\mathbf{y})$ can be computed as

$$\begin{aligned} \text{Cov}_{\mathbf{y}|\theta} (\Lambda(\mathbf{y}, \theta), \hat{\theta}(\mathbf{y})) &= \mathbb{E}_{\mathbf{y}|\theta} [(\Lambda(\mathbf{y}, \theta) - \mathbb{E}[\Lambda(\mathbf{y}, \theta)]) (\hat{\theta}(\mathbf{y}) - \mathbb{E}[\hat{\theta}(\mathbf{y})])] \\ &= \mathbb{E}_{\mathbf{y}|\theta} [\Lambda(\mathbf{y}, \theta) \hat{\theta}(\mathbf{y})] - \mathbb{E}_{\mathbf{y}|\theta} [\Lambda(\mathbf{y}, \theta)] \mathbb{E}_{\mathbf{y}|\theta} [\hat{\theta}(\mathbf{y})] \\ &= \mathbb{E}_{\mathbf{y}|\theta} [\Lambda(\mathbf{y}, \theta) \hat{\theta}(\mathbf{y})] \\ &= \int p(\mathbf{y} | \theta) \Lambda(\mathbf{y}, \theta) \hat{\theta}(\mathbf{y}) d\mathbf{y} \\ &= \int \frac{\partial}{\partial \theta} p(\mathbf{y} | \theta) \hat{\theta}(\mathbf{y}) d\mathbf{y} \\ &= \frac{\partial}{\partial \theta} \int p(\mathbf{y} | \theta) \hat{\theta}(\mathbf{y}) d\mathbf{y} \\ &= \frac{\partial}{\partial \theta} (\mathbb{E}_{\mathbf{y}|\theta} [\hat{\theta}(\mathbf{y})] - \theta) + 1 \\ &= \frac{\partial}{\partial \theta} b_{\hat{\theta}} + 1. \end{aligned}$$

Using the Cauchy-Schwarz inequality, we have the following bound,

$$\begin{aligned} \text{Cov}_{\mathbf{y}|\theta} (\Lambda(\mathbf{y}, \theta), \hat{\theta}(\mathbf{y}))^2 &= \left(\mathbb{E}_{\mathbf{y}|\theta} [\Lambda(\mathbf{y}, \theta) (\hat{\theta}(\mathbf{y}) - \mathbb{E}[\hat{\theta}(\mathbf{y})])] \right)^2 \\ \left(\frac{\partial}{\partial \theta} b_{\hat{\theta}} + 1 \right)^2 &\leq \mathbb{E}_{\mathbf{y}|\theta} [\Lambda(\mathbf{y}, \theta)^2] \cdot \mathbb{E}_{\mathbf{y}|\theta} [(\hat{\theta}(\mathbf{y}) - \mathbb{E}[\hat{\theta}(\mathbf{y})])^2] \\ &= \mathcal{I}_n(\theta) \cdot \mathbb{E}_{\mathbf{y}|\theta} [(\hat{\theta}(\mathbf{y}) - \theta - \mathbb{E}[\hat{\theta}(\mathbf{y}) - \theta])^2] \\ &= \mathcal{I}_n(\theta) \cdot \left(\mathbb{E}_{\mathbf{y}|\theta} [(\hat{\theta}(\mathbf{y}) - \theta)^2] + (\mathbb{E}_{\mathbf{y}|\theta} [\hat{\theta}] - \theta)^2 \right. \\ &\quad \left. - 2\mathbb{E}_{\mathbf{y}|\theta} [\hat{\theta}(\mathbf{y}) - \theta] \right) \\ &\leq \mathcal{I}_n(\theta) \cdot \left(\mathbb{E}_{\mathbf{y}|\theta} [(\hat{\theta}(\mathbf{y}) - \theta)^2] - b_{\hat{\theta}}^2 \right). \end{aligned}$$

Re-arranging yields

$$\mathbb{E}_{\mathbf{y}|\theta} [(\hat{\theta}(\mathbf{y}) - \theta)^2] \geq \frac{\left(\frac{\partial}{\partial \theta} b_{\hat{\theta}} + 1 \right)^2}{\mathcal{I}_n(\theta)} + b_{\hat{\theta}}^2.$$

■

Note the trade-off between variance and bias. If $\frac{\partial}{\partial \theta} b_{\hat{\theta}} < 0$, then the variance might be smaller than the variance of the best unbiased estimator.

Corollary. Let $\hat{\theta}$ be an unbiased estimator—i.e., $b_{\hat{\theta}} = 0$ —then

$$\text{Var}[\hat{\theta}(\mathbf{y})] \geq \frac{1}{\mathcal{I}_n(\theta)}.$$

Lemma 3.2. The maximum likelihood estimator $\hat{\theta}_{\text{ML}}$ is asymptotically efficient,

$$\lim_{n \rightarrow \infty} \text{Var}[\hat{\theta}_{\text{ML}}] = \frac{1}{\mathcal{I}_n(\theta)}.$$

However, the MLE estimator is not necessarily efficient for finite samples.

4 Regression

In regression, we try to estimate a function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ that best fits a given dataset $\{(x_i, y_i)\}_{i=1}^n \subseteq \mathbb{R}^d \times \mathbb{R}$. I.e., we try to minimize the following loss function,

$$\ell(\beta) = \frac{1}{n} \sum_{i=1}^n \|f(x_i) - y_i\|^2.$$

4.1 Linear regression

In linear regression, we assume the underlying data to be linear with a fixed noise,

$$Y | X = x \sim \mathcal{N}(\beta_*^\top x, \sigma^2)$$

for some ground truth β_* . We parametrize f as a linear function,

$$f(x; \beta) = \beta^\top x.$$

Under these assumptions, the minimizer of the loss function is analytically tractable—the ordinary least squares estimator (OLSE),

$$\hat{\beta} = (X^\top X)^{-1} X^\top y,$$

where $X \in \mathbb{R}^{d \times n}$ is the design matrix with respective outputs $y \in \mathbb{R}^n$.

In practice, it is important to remove outliers, because linear models can be heavily influenced by them. Also, the data should be standardized, such that all features are on the same scale, because differences in scale can make matrix inversion unstable.

TODO: Curse of dimensionality [Sur and Candès, 2019], Figures 4.1 and 4.2.

Further, if features are collinear, then the model might only learn the correlation with the target variable of one of them and discard the other. In addition, some singular values will equal zero. This makes matrix inversion unstable. This can easily be shown by considering the singular value decomposition $X = UDV^\top$, then

$$\begin{aligned} \hat{\beta} &= (X^\top X)^{-1} X^\top y \\ &= (VD^\top U^\top UDV^\top)^{-1} VD^\top U^\top y \\ &= (VD^\top DV^\top)^{-1} VD^\top U^\top y \\ &= V(D^{-1})^\top D^{-1} V^\top VD^\top U^\top y \\ &= V(D^{-1})^\top D^{-1} D^\top U^\top y \\ &= VD^{-1} U^\top y. \end{aligned}$$

The inversion of D is unstable if the singular values are small. The solution to this is to remove collinear features and/or to add regularization.




Figure 4.1. β^* is shown as black marks and $\hat{\beta}$ is indicated by the marks. As can be seen, $\hat{\beta}$ is overestimated for indices where $\beta_i^* \neq 0$.




Figure 4.2. True *vs.* predicted probability—the estimator is overconfident in its predictions.

Consider the mean-squared error loss. The risk (expected loss) can be decomposed as follows into bias, variance, and noise terms,

$$\mathbb{E}[(f(\mathbf{x}) - y)^2] = \underbrace{(\mathbb{E}[f(\mathbf{x})] - \mathbb{E}[y])^2}_{\text{squared bias}} + \underbrace{\text{Var}[f(\mathbf{x})]}_{\text{variance}} + \underbrace{\mathbb{E}[(\mathbb{E}[y] - y)^2]}_{\text{noise}}.$$

The OLSE is the minimum variance unbiased estimator.² However, this does not mean it is the best, because introducing some bias may considerably decrease the variance. Bayesianism adds bias by introducing a prior—priors often induce a regularizing term.³

² This is proven by the Gauss-Markov theorem.

³ E.g., a Gaussian prior induces an ℓ_2 -norm regularizing term.

4.2 Polynomial regression

In polynomial regression, we construct a feature function of all possible polynomials, e.g.,

$$\phi([x_1, x_2]) = [1, x_1, x_2, x_1^2, x_1x_2, x_2^2, x_1^3, x_1^2x_2, \dots].$$

We then perform linear regression in this space,

$$\psi(\mathbf{x}; \boldsymbol{\beta}) = \langle \boldsymbol{\beta}, \phi(\mathbf{x}) \rangle.$$

The problem is that this space is infinitely dimensional, and the inner product is ill-defined in an infinitely dimensional space.⁴ We can solve this by transforming the vector such that inner products cannot diverge by introducing a data-dependent scalar,

$$\phi(\mathbf{x}) = \alpha(\mathbf{x}) \cdot \tilde{\phi}(\mathbf{x}), \quad \alpha(\mathbf{x}) > 0.$$

⁴ This is due to $\langle \phi(\mathbf{x}), \phi(\mathbf{x}') \rangle = \infty$ for some \mathbf{x}, \mathbf{x}' , e.g., $\mathbf{x} = \mathbf{x}' = [1, 1]$.

$\tilde{\phi}$ contains all polynomials in this equation.

An inner product w.r.t. this feature function is a valid inner product. Such transformations can have a closed form for the inner product, called kernelization. Commonly, the radial basis function (RBF) kernel is used,

$$\langle \phi(\mathbf{x}), \phi(\mathbf{x}') \rangle = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\sigma^2}\right), \quad \sigma \in \mathbb{R}.$$

Let $\Phi \in \mathbb{R}^{n \times \infty}$ contain all feature vectors, then the OLSE for polynomial regression can be computed by

$$\hat{\boldsymbol{\beta}} = (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y}, \quad \Phi \in \mathbb{R}^{n \times \infty}.$$

However, $\Phi^\top \Phi$ cannot be computed, because it is $\infty \times \infty$ -dimensional. We can fix this by observing that $\Phi \Phi^\top \in \mathbb{R}^{n \times n}$ and rewriting the OLSE as

$$\begin{aligned} \hat{\boldsymbol{\beta}} &= (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y} \\ &= (\Phi^\top \Phi)^{-1} \Phi^\top (\Phi \Phi^\top)^{-1} (\Phi \Phi^\top) \Phi^\top \mathbf{y} \\ &= (\Phi^\top \Phi)^{-1} \Phi^\top \Phi \Phi^\top (\Phi \Phi^\top)^{-1} \mathbf{y} \\ &= \Phi^\top (\Phi \Phi^\top)^{-1} \mathbf{y}. \end{aligned}$$

Let $\mathbf{K} = \Phi\Phi^\top$, then it only contains kernel evaluations— $k_{ij} = \langle\phi(\mathbf{x}_i), \phi(\mathbf{x}_j)\rangle$. The next problem is that Φ^\top is still infinitely dimensional. However, this is not a problem, since when we want to make a prediction, we do the following,

$$\begin{aligned}\psi(\hat{\mathbf{x}}) &= \langle\phi(\hat{\mathbf{x}}), \hat{\boldsymbol{\beta}}\rangle \\ &= \phi(\hat{\mathbf{x}})^\top \Phi^\top \mathbf{K}^{-1} \mathbf{y}.\end{aligned}$$

Let $\mathbf{k}(\hat{\mathbf{x}}) = \phi(\hat{\mathbf{x}})^\top \Phi^\top$, then it only contains kernel evaluations with the new point— $k_i(\hat{\mathbf{x}}) = \langle\phi(\hat{\mathbf{x}}), \phi(\mathbf{x}_i)\rangle$. In conclusion, we can make predictions by

$$\psi(\hat{\mathbf{x}}) = \mathbf{k}(\hat{\mathbf{x}}) \mathbf{K}^{-1} \mathbf{y}, \quad \mathbf{k}(\hat{\mathbf{x}}) \in \mathbb{R}^{1 \times n}, \mathbf{K} \in \mathbb{R}^{n \times n}, \mathbf{y} \in \mathbb{R}^n.$$

However, the problem with this approach is that it takes $O(n^3)$ to make a prediction—it scales in the amount of data points.

4.3 Regularization

TODO: Ridge regression— ℓ_2 -norm.

TODO: LASSO— ℓ_1 -norm induces sparsity. Better interpretability properties.

TODO: Figure showing coefficient weights under different regularization factors.

5 Causality

In general, models do not distinguish between causal and non-causal factors in the feature space. Therefore, they might identify non-causal factors as highly correlating with the output variable. *E.g.*, when classifying images as either depicting a cow or a camel, the model might identify the background as an important feature, because cows are usually in grass and camels are usually in the desert. However, if this classifier were to see a cow with a different background, it would fail.

The following are examples of causal fallacies, where one might conclude that X causes Y ,

- Reverse causality, where Y actually predicts X and not the other way around;
- Third-cause fallacy, where there is a confounding factor Z that influences both X and Y ;
- Bidirectional causation, where X influence each other;

A domain shift happens when the test samples are drawn from a different distribution than the training samples. *E.g.*, COVID-19 detection models trained on a western population might not perform well on an eastern population.

Shortcut learning happens when there is a spurious correlation between causal and non-causal features in the training dataset. The resulting estimator abuses the non-causal features to generalize in the testing dataset. The solution to this is to encode the features such that they do not depend on the environment.

5.1 Counterfactual invariance

Let X be the feature vector representing the object and let Y be a target variable of interest, and let f be the function that estimates Y from X . Further, let W describe features that influence X , but should not influence the estimator f . Let a counterfactual be denoted as $X(w)$, which is the feature vector we would have obtained if we would have had $W = w$, after the fact. Then, the estimator f is counterfactually invariant if

$$f(X(w)) = f(X(w')), \quad \forall w, w' \in \text{range}(W).$$

In words, f is not influenced by the value of W . One way of obtaining a counterfactually invariant estimator is by extending the training dataset to contain enough counterfactuals. This can, for example, be achieved by data augmentation. However, this is not always possible.

Let X_A be the parts of X causally influenced by A , and let A^\perp be the set of variables independent of A . Then, f is counterfactually invariant if and only if f only depends on X_{W^\perp} .




Figure 5.1. Causal graph.




Figure 5.2. Anti-causal graph.

Theorem 5.1 (Necessary conditions for counterfactual invariance [Veitch et al., 2021]). If f is a counterfactually invariant predictor, then

- In the anti-causal scenario, $f(\mathbf{X}) \perp \mathbf{W} \mid Y$;
- In the causal scenario without selection (but possibly confounded), $f(\mathbf{X}) \perp \mathbf{W}$;
- In the causal scenario without confoundedness (but possibly with selection), $f(\mathbf{X}) \perp \mathbf{W} \mid Y$, as long as $\mathbf{X} \perp Y \mid \mathbf{X}_{\mathbf{W}^\perp}, \mathbf{W}$.

Proof. This can be proven by using d-separation. ■

6 Gaussian processes

Let the inputs be $\mathbf{X} \in \mathbb{R}^{n \times d}$, the outputs be $\mathbf{y} \in \mathbb{R}^n$, and the weight matrix be $\boldsymbol{\beta} \in \mathbb{R}^d$, then linear regression models the generative process as

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}).$$

This is equivalent to defining a Gaussian over the predictions,

$$\mathbf{y} \mid \mathbf{X}, \boldsymbol{\beta} \sim \mathcal{N}(\mathbf{X}\boldsymbol{\beta}, \sigma^2 \mathbf{I}).$$

BLR (*Bayesian Linear Regression*) extends linear regression by defining a prior over the regression coefficients,

$$\boldsymbol{\beta} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Lambda}^{-1}), \quad \boldsymbol{\Lambda} \in \mathbb{R}^{d \times d}.$$

The posterior can be analytically computed as

$$\boldsymbol{\beta} \mid \mathbf{X}, \mathbf{y} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma}),$$

where

$$\boldsymbol{\mu} = \frac{1}{\sigma^2} \boldsymbol{\Sigma} \mathbf{X}^\top \mathbf{y}, \quad \boldsymbol{\Sigma} = \sigma^2 (\mathbf{X}^\top \mathbf{X} + \sigma^2 \boldsymbol{\Lambda})^{-1}.$$

Proof.

$$\begin{aligned} p(\boldsymbol{\beta} \mid \mathbf{X}, \mathbf{y}) &\propto p(\mathbf{y} \mid \mathbf{X}, \boldsymbol{\beta}) \cdot p(\boldsymbol{\beta}) && \text{Bayes' rule.} \\ &= \mathcal{N}(\mathbf{y}; \mathbf{X}\boldsymbol{\beta}, \sigma^2 \mathbf{I}) \cdot \mathcal{N}(\boldsymbol{\beta}; \mathbf{0}, \boldsymbol{\Lambda}^{-1}) \\ &\propto \exp\left(-\frac{1}{2} \left(\frac{1}{\sigma^2} \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|^2 + \boldsymbol{\beta}^\top \boldsymbol{\Lambda} \boldsymbol{\beta} \right)\right) \\ &= \exp\left(-\frac{1}{2} \left(\frac{1}{\sigma^2} (\|\mathbf{y}\|^2 + \|\mathbf{X}\boldsymbol{\beta}\|^2 - 2\mathbf{y}^\top \mathbf{X}\boldsymbol{\beta}) + \boldsymbol{\beta}^\top \boldsymbol{\Lambda} \boldsymbol{\beta} \right)\right) && \text{Cosine theorem.} \\ &= \exp\left(-\frac{1}{2} \left(\frac{1}{\sigma^2} \mathbf{y}^\top \mathbf{y} + \frac{1}{\sigma^2} \boldsymbol{\beta}^\top \mathbf{X}^\top \mathbf{X} \boldsymbol{\beta} - \frac{2}{\sigma^2} \mathbf{y}^\top \mathbf{X} \boldsymbol{\beta} + \boldsymbol{\beta}^\top \boldsymbol{\Lambda} \boldsymbol{\beta} \right)\right) \\ &= \exp\left(-\frac{1}{2} \left(\boldsymbol{\beta}^\top \left(\frac{1}{\sigma^2} \mathbf{X}^\top \mathbf{X} + \boldsymbol{\Lambda} \right) \boldsymbol{\beta} - \frac{2}{\sigma^2} \mathbf{y}^\top \mathbf{X} \boldsymbol{\beta} + \frac{1}{\sigma^2} \mathbf{y}^\top \mathbf{y} \right)\right) \\ &= \exp\left(-\frac{1}{2} \left(\boldsymbol{\beta}^\top \boldsymbol{\Sigma}^{-1} \boldsymbol{\beta} - \frac{1}{\sigma^2} \boldsymbol{\beta}^\top \mathbf{X}^\top \mathbf{y} - \frac{1}{\sigma^2} \mathbf{y}^\top \mathbf{X} \boldsymbol{\beta} + \frac{1}{\sigma^2} \mathbf{y}^\top \mathbf{y} \right)\right) \\ &= \exp\left(-\frac{1}{2} \left(\boldsymbol{\beta}^\top \boldsymbol{\Sigma}^{-1} \left(\boldsymbol{\beta} - \frac{1}{\sigma^2} \boldsymbol{\Sigma} \mathbf{X}^\top \mathbf{y} \right) - \frac{1}{\sigma^2} \mathbf{y}^\top \mathbf{X} \left(\boldsymbol{\beta} - \frac{1}{\sigma^2} \boldsymbol{\Sigma} \mathbf{X}^\top \mathbf{y} \right) \right)\right) \\ &= \exp\left(-\frac{1}{2} \left(\left(\boldsymbol{\beta} - \frac{1}{\sigma^2} \boldsymbol{\Sigma} \mathbf{X}^\top \mathbf{y} \right)^\top \boldsymbol{\Sigma}^{-1} \left(\boldsymbol{\beta} - \frac{1}{\sigma^2} \boldsymbol{\Sigma} \mathbf{X}^\top \mathbf{y} \right) \right)\right) \\ &\propto \mathcal{N}\left(\boldsymbol{\beta}; \frac{1}{\sigma^2} \boldsymbol{\Sigma} \mathbf{X}^\top \mathbf{y}, \boldsymbol{\Sigma}\right). \end{aligned}$$

■

In general, we do not know the true weights that generated the data points. But, we can now define a joint distribution over output variables with unknown weights,

$$\mathbf{y} \mid \mathbf{X} \sim \mathcal{N}(\mathbf{0}, \mathbf{X} \boldsymbol{\Lambda}^{-1} \mathbf{X}^\top + \sigma^2 \mathbf{I}).$$

Proof. As we saw earlier, we compute outputs as follows,

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}).$$

However, now the weights are unknown, so we make use of its prior to compute the distribution over \mathbf{y} ,

$$\begin{aligned} \mathbb{E}[\mathbf{y}] &= \mathbb{E}[\mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}] \\ &= \mathbf{X}\mathbb{E}[\boldsymbol{\beta}] + \mathbb{E}[\boldsymbol{\epsilon}] \\ &= \mathbf{0}. \end{aligned}$$

$$\begin{aligned} \text{Cov}[\mathbf{y}] &= \mathbb{E}[(\mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon})(\mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon})^\top] \\ &= \mathbb{E}[\mathbf{X}\boldsymbol{\beta}\boldsymbol{\beta}^\top \mathbf{X}^\top] + \mathbb{E}[\mathbf{X}\boldsymbol{\beta}\boldsymbol{\epsilon}^\top] + \mathbb{E}[\boldsymbol{\epsilon}\boldsymbol{\beta}^\top \mathbf{X}^\top] + \mathbb{E}[\boldsymbol{\epsilon}\boldsymbol{\epsilon}^\top] \\ &= \mathbf{X}\mathbb{E}[\boldsymbol{\beta}\boldsymbol{\beta}^\top] \mathbf{X}^\top + \mathbf{X}\mathbb{E}[\boldsymbol{\beta}]\mathbb{E}[\boldsymbol{\epsilon}^\top] + \mathbb{E}[\boldsymbol{\epsilon}]\mathbb{E}[\boldsymbol{\beta}^\top] \mathbf{X}^\top + \sigma^2 \mathbf{I} \\ &= \mathbf{X}\boldsymbol{\Lambda}^{-1} \mathbf{X}^\top + \sigma^2 \mathbf{I}. \end{aligned}$$

$\boldsymbol{\beta}$ and $\boldsymbol{\epsilon}$ are independent.

■

GPs (*Gaussian Processes*) generalize BLR by observing that we can kernelize the covariance matrix,

$$k(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^\top \boldsymbol{\Lambda}^{-1} \mathbf{x}_j.$$

We could instead use any other kernel function to model other functions.

6.1 Kernels

Kernel functions specify the similarity between any two data points. They encode assumptions about the function that is to be learned.

Definition 6.1 (Kernel function). A kernel function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ over $\mathcal{X} \subset \mathbb{R}^d$ must satisfy the following properties,

$$\begin{aligned} k(\mathbf{x}, \mathbf{x}') &= k(\mathbf{x}', \mathbf{x}) \\ \int k(\mathbf{x}, \mathbf{x}') f(\mathbf{x}) f(\mathbf{x}') d\mathbf{x} d\mathbf{x}' &\geq 0, \quad \forall f \in L_2. \end{aligned}$$

Symmetry.

Positive semi-definiteness.

Definition 6.2 (Stationary and isotropic). A kernel $k(\mathbf{x}, \mathbf{x}')$ is stationary if it only depends on $\mathbf{x} - \mathbf{x}'$. Further, it is isotropic if it only depends on $\|\mathbf{x} - \mathbf{x}'\|_2$.

The following are common kernels,

$$k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^\top \mathbf{x}'$$

Linear kernel.

$$k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}' + 1)^p, \quad p \in \mathbb{N}$$

Polynomial kernel.

$$k(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|_2^2}{\ell^2}\right), \quad \ell \in \mathbb{R}$$

RBF (*Radial Basis Function*) kernel.

$$k(\mathbf{x}, \mathbf{x}') = \tanh(\kappa \mathbf{x}^\top \mathbf{x}') - b, \quad \kappa, b \in \mathbb{R}$$

Sigmoid kernel.

Different kernels have different invariance properties, such as invariance to rotation or translation. In order to learn invariances from data, many samples are needed. So, it might be better to encode them if we know them a priori.

Given two kernel functions k_1, k_2 defined on the same data space and positive scalar $c > 0$, the following are also valid kernels,

$$k(x, x') = k_1(x, x') + k_2(x, x')$$

$$k(x, x') = k_1(x, x') \cdot k_2(x, x')$$

$$k(x, x') = c \cdot k_1(x, x')$$

$$k(x, x') = \exp(k_1(x, x')).$$

In practice, the kernels are often composed together and hyperparameters are determined by performance on a held out validation dataset.

6.2 Prediction

As we saw earlier in the case of BLR, the marginal is jointly Gaussian,

$$\begin{bmatrix} y \\ y^* \end{bmatrix} \mid X, x^* \sim \mathcal{N}\left(\mathbf{0}, \begin{bmatrix} K + \sigma^2 I & k(X, x^*) \\ k(x^*, X) & k(x^*, x^*) \end{bmatrix}\right).$$

Theorem 6.3 (Conditional Gaussian distribution). Given

$$\begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix}, \begin{bmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{bmatrix}\right).$$

Then, the conditional Gaussian is given by

$$x_2 \mid x_1 = z \sim \mathcal{N}\left(\mu_2 + \Sigma_{21}\Sigma_{11}^{-1}(z - \mu_1), \Sigma_{22} - \Sigma_{21}\Sigma_{11}^{-1}\Sigma_{12}\right).$$

Using the above theorem, we can easily compute the predictive distribution,

$$y^* \mid x^*, X, y \sim \mathcal{N}\left(k^\top (K + \sigma^2 I)^{-1} y, c - k^\top (K + \sigma^2 I)^{-1} k\right),$$

where

$$K = K(X, X), \quad k = k(X, x^*), \quad c = k(x^*, x^*).$$

Now, we can compute a prediction along with its uncertainty. However, the problem with this approach is that it has $\mathcal{O}(n^3)$ runtime.




Figure 6.1. Fitted Gaussian process (RBF kernel, $\ell = 1.86$) with its 95% confidence interval.

7 Uncertainty quantification

7.1 Statistical model validation

In statistical modeling, we typically want to predict a target random variable $Y \in \mathcal{Y}$ given a random vector $\mathbf{X} \in \mathcal{X}$. Formally, we want to find a function $f : \mathcal{X} \rightarrow \mathcal{Y}$ that minimizes the expected risk,

$$\mathcal{R}(f) \doteq \mathbb{E}_{\mathbf{X}, Y}[\mathbb{1}\{f(\mathbf{X}) \neq Y\}].$$

However, this function is intractable, because (1) we do not have access to the joint distribution over \mathbf{X}, Y , (2) it is intractable to find f without any assumptions on its structure, and (3) it is unclear how to minimize the expectation of the indicator function. To resolve (1), we obtain a dataset of samples from the joint distribution. Problem (2) can be resolved by restricting the functions to a parameterized hypothesis space \mathcal{H} . Lastly, the solution to (3) is to approximate the indicator function by a differentiable loss function $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$. As a result, given a dataset $\mathcal{Z} = \{(x_i, y_i)\}_{i=1}^n$, we can approximate the expected risk by the empirical risk,

$$\hat{\mathcal{R}}(f) \doteq \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)).$$

We denote the minimizer of this function by \hat{f} .

Since \hat{f} is obtained by training on the dataset \mathcal{Z} , we need to find a way to measure the performance of \hat{f} on unseen data points—we want to estimate $\mathcal{R}(\hat{f})$. Furthermore, let $\mathcal{A} : \mathcal{P}(\mathcal{X} \times \mathcal{Y}) \rightarrow \mathcal{H}$ be the algorithm that maps datasets to functions, then we want to estimate its expected risk,

$$\mathcal{R}(\mathcal{A}) \doteq \mathbb{E}_{\mathcal{Z}}[\mathcal{R}(\mathcal{A}(\mathcal{Z}))].$$

Cross-validation. Cross-validation works as follows,

1. Partition the data \mathcal{Z} in K equally sized disjoint subsets, such that

$$\mathcal{Z} = \bigcup_{k=1}^K \mathcal{Z}_k;$$

2. Use \mathcal{A} to produce K estimators \hat{f}^{-k} from $\mathcal{Z} - \mathcal{Z}_k$ for all $k \in [K]$;
3. Estimate the expected risk by

$$\mathcal{R}^{\text{CV}}(\mathcal{A}) \doteq \frac{1}{n} \sum_{i=1}^n \ell(y_i, \hat{f}^{-(k(i))}(x_i)),$$

where $k : [n] \rightarrow [K]$ maps i to the partition such that $x_i \in \mathcal{Z}_{k(i)}$.

In conclusion, the loss terms in the expected risk depend on approximators that have not seen the data it is evaluating.

Bootstrap. The bootstrap method is used for measuring the distribution over statistical parameters. It does so by creating B bootstrap samples by sampling with replacement from the dataset. For each bootstrap sample, the parameters are computed. Using these parameters, we can estimate the variability and uncertainty of the parameter. Note that this works for any type of parameter that can be estimated from the dataset—not only model parameters. Let $S : \mathcal{P}(\mathcal{X} \times \mathcal{Y}) \rightarrow \mathbb{R}^d$ be a function that processes datasets and outputs parameters, then we can compute the statistics of this parameter as follows,

1. Draw B samples \mathcal{Z}^{*b} of size n from \mathcal{Z} with replacement;
2. Compute an estimate $S(\mathcal{Z}^{*b})$ for each bootstrap sample.

Now, we can use this set of parameter estimates to compute its statistics. For example, we can compute its mean and covariance,

$$\begin{aligned}\mu(S) &= \frac{1}{B} \sum_{b=1}^B S(\mathcal{Z}^{*b}) \\ \Sigma(S) &= \frac{1}{B-1} \sum_{b=1}^B \left(S(\mathcal{Z}^{*b}) - \mu(S) \right) \left(S(\mathcal{Z}^{*b}) - \mu(S) \right)^\top.\end{aligned}$$

The algorithm \mathcal{A} is a function that processes a dataset and outputs an estimator, which is used to compute a loss term. We can use this to estimate the empirical risk across n data points,

$$\hat{\mathcal{R}}^{\text{BS}}(\mathcal{A}) \doteq \frac{1}{n \cdot B} \sum_{b=1}^B \sum_{i=1}^n \ell(y_i, f^{*b}(x_i)).$$

However, computing this validation metric includes data that the bootstrap estimates were trained on—the estimate is overly optimistic. Thus, we can use the bootstrap method to estimate the expected risk $\mathcal{R}(\mathcal{A})$ as follows,

$$\mathcal{R}^{\text{BS}}(\mathcal{A}) \doteq \frac{1}{n} \sum_{i=1}^n \frac{1}{|\mathcal{C}^{-i}|} \sum_{b \in \mathcal{C}^{-i}} \ell(y_i, \hat{f}^{*b}(x_i)),$$

where \mathcal{C}^{-i} contains the indices of the bootstraps that do not contain observation (x_i, y_i) .

In order to correct for the optimism of $\hat{\mathcal{R}}(\mathcal{A})^{\text{BS}}$, we combine it with \mathcal{R}^{BS} ,

$$\mathcal{R}^{(0.632)} \doteq 0.368 \hat{\mathcal{R}}^{\text{BS}} + 0.632 \mathcal{R}^{\text{BS}}.$$

Here, a weight of 0.632 is used because that is the probability that a sample (x_i, y_i) appears at least once in a bootstrap sample of size n ,

$$1 - \left(1 - \frac{1}{n}\right)^n \rightarrow 1 - \frac{1}{e} = 0.632.$$

TODO: Figure out why we are weighting it, because the optimism is already removed in \mathcal{R}^{BS} .

7.2 Uncertainty in linear models

Suppose we have n observations in our dataset $\{(x_i, y_i)\}_{i=1}^n$. Let $X \in \mathbb{R}^{n \times d}$ and $\mathbf{y} \in \mathbb{R}^n$ be the design matrix and output vector. We further assume

$$\mathbf{y} \mid X \sim \mathcal{N}(X\boldsymbol{\beta}^*, \sigma^2 \mathbf{I}).$$

As we have seen before, the ordinary least squares estimator is computed as follows,

$$\hat{\boldsymbol{\beta}} = (X^\top X)^{-1} X^\top \mathbf{y}.$$

Thus, we have the following distribution over estimators,

$$\hat{\boldsymbol{\beta}} \sim \mathcal{N}(\boldsymbol{\beta}^*, \sigma^2 (X^\top X)^{-1}).$$

An unbiased estimator of σ^2 is

$$\hat{\sigma}^2 = \frac{1}{n-d} \sum_{i=1}^n (\hat{\boldsymbol{\beta}}^\top \mathbf{x}_i - y_i)^2.$$

Thus, we can approximate a $1 - \alpha$ confidence interval for β_j^* by

$$\hat{\beta}_j \pm z_{\alpha/2} \hat{\epsilon}(\hat{\beta}_j),$$

where $z_{\alpha/2} = \Phi^{-1}(\alpha/2)$, Φ is the CDF of the standard Gaussian, and $\epsilon(\hat{\beta}_j)$ is the j -th diagonal element of $\sigma^2 (X^\top X)^{-1}$

TODO: Find out how the $1 - \alpha$ confidence interval of β_j^* is computed.

7.3 Statistical testing

Assume we have a hypothesis set $\mathcal{H} = \{p(\cdot \mid \boldsymbol{\theta}) \mid \boldsymbol{\theta} \in \Theta\}$ and let $\boldsymbol{\theta}^* \in \Theta$ be the (unknown) true parameter. Furthermore, we are given a null hypothesis $H_0 : \boldsymbol{\theta}^* \in \Theta_0$ and an alternative hypothesis $H_1 : \boldsymbol{\theta}^* \in \Theta_1$ for predefined $\Theta_0, \Theta_1 \subseteq \Theta$. Lastly, we are given n samples $x_1, \dots, x_n \sim p(\cdot \mid \boldsymbol{\theta}^*)$ and a test statistic $t : \mathcal{X}^n \rightarrow \mathbb{R}$. The goal is to find a critical value $c \in \mathbb{R}$ for the test t such that

$$\mathbb{P}(t(X_1, \dots, X_n) \geq c \mid \boldsymbol{\theta})$$

is “low” when $\boldsymbol{\theta} \in \Theta_0$ and “high” when $\boldsymbol{\theta} \in \Theta_1$. So, we accept the null hypothesis when the probability of the test statistic being greater than c is low and we reject it if it is high.

We want to minimize the probability of incorrectly choosing H_1 as this is worse than incorrectly choosing H_0 . We can quantify this notion of risk as

$$\alpha_c \doteq \sup_{\boldsymbol{\theta} \in \Theta_0} \mathbb{P}(t(x_1, \dots, x_n) \geq c \mid \boldsymbol{\theta}).$$

Intuitively, this is the maximum probability of incorrectly choosing H_1 , because the test is passed, but the null hypothesis holds. Note that $\alpha_c \rightarrow 0$

as $c \rightarrow \infty$. So, $c^* \rightarrow \infty$ minimizes the risk. But, this comes with the problem that we would never accept H_1 .

The solution to this problem is to forget about choosing the optimal critical value a priori and running an experiment to obtain a realization x_1, \dots, x_n . We then run the test with the realization $t(x_1, \dots, x_n)$ and compute the risk of the least risky critical value that would incorrectly reject H_0 ,

$$p = \inf_{c \in \mathbb{R}} \{\alpha_c \mid t(x_1, \dots, x_n) \geq c\}.$$

One can show that this is the common p -value,

$$p \doteq \sup_{\theta \in \Theta_0} \mathbb{P}(t(X_1, \dots, X_n) \geq t(x_1, \dots, x_n) \mid \theta).$$

Intuitively, this is the probability that the test is higher than our current observation if we run the experiment again.

Wald test. The Wald test is an example of a statistical test. Let $\hat{\theta}$ be an estimator of θ with standard deviation $\hat{\sigma}$ that is asymptotically normal. Let the null hypothesis H_0 be $\theta = \theta_0$ and alternative hypothesis H_1 be $\theta \neq \theta_0$, then the Wald test statistic is

$$W = \frac{(\hat{\theta} - \theta_0)^2}{\hat{\sigma}^2}.$$

TODO: Figure out what the point of this is in the context of machine learning. Also, figure out how this can be done practically, because I have no idea how one would compute α_c or p .

7.4 Bayesian neural networks

As we have seen, BLR incorporates Bayesian uncertainty into linear regression. In a similar fashion, BNNs (*Bayesian Neural Networks*) incorporate Bayesian inference into neural networks.

Neural networks. We define a neural network as having L alternating linear layers and element-wise activation functions,

$$\begin{aligned} \mathbf{z}_0 &= \mathbf{x} \\ \mathbf{z}_\ell &= \phi(\mathbf{a}_\ell) \\ \mathbf{a}_\ell &= \mathbf{W}_\ell \mathbf{z}_{\ell-1} + \mathbf{b}_\ell, \end{aligned}$$

where ϕ is a differentiable activation function and \mathbf{a} is its pre-activation. The output of the neural network is then

$$f(\mathbf{x} \mid \boldsymbol{\theta}) = \mathbf{z}_L, \quad \boldsymbol{\theta} = \{\mathbf{W}_\ell, \mathbf{b}_\ell\}_{\ell=1}^L.$$

In practice, the weights of the neural network are chosen to minimize the empirical risk,

$$\hat{\boldsymbol{\theta}} \in \operatorname{argmin}_{\boldsymbol{\theta} \in \Theta} \hat{\mathcal{R}}(f(\cdot \mid \boldsymbol{\theta})) \doteq \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(\mathbf{x}_i \mid \boldsymbol{\theta})).$$

In general, the weights are optimized by gradient descent,

$$\boldsymbol{\theta}_t = \boldsymbol{\theta}_{t-1} - \eta_t \nabla_{\boldsymbol{\theta}} \hat{\mathcal{R}}(f(\cdot | \boldsymbol{\theta}_{t-1})).$$

For large datasets, it is too expensive to compute the gradient, so we instead use SGD (*Stochastic Gradient Descent*), which replaces the gradient by an unbiased estimate,

$$\boldsymbol{\theta}_t = \boldsymbol{\theta}_{t-1} - \eta_t \nabla_{\boldsymbol{\theta}} \left(\frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \ell(y_i, f(\mathbf{x}_i | \boldsymbol{\theta}_{t-1})) \right), \quad \mathcal{B} \subseteq [n],$$

where \mathcal{B} contains uniformly sampled sample indices.

Bayesian. The disadvantage of (stochastic) gradient descent is that it yields only a single point estimate of the weights. There is no quantification of uncertainty in this estimate, which leads to overconfidence problems that can result in poor generalization in the presence of domain shifts. We will use a Bayesian approach to alleviate this problem.

We must first define a prior $p(\boldsymbol{\theta})$ over the weights $\boldsymbol{\theta} \in \Theta$,

$$\boldsymbol{\theta} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}).$$

Then, we can define the likelihood function $p(\mathcal{D} | \boldsymbol{\theta})$ that computes the likelihood of the dataset \mathcal{Z} for a given $\boldsymbol{\theta}$,

$$p(\mathcal{Z} | \boldsymbol{\theta}) = \prod_{(\mathbf{x}, y) \in \mathcal{Z}} p(y | \mathbf{x}, \boldsymbol{\theta}),$$

where $p(y | \mathbf{x}, \boldsymbol{\theta})$ is a probability computed by the neural network. Lastly, we use Bayes rule to compute the posterior over weights,

$$p(\boldsymbol{\theta} | \mathcal{Z}) = \frac{p(\boldsymbol{\theta}) p(\mathcal{Z} | \boldsymbol{\theta})}{p(\mathcal{Z})}.$$

The problem is that $p(\mathcal{Z}) = \int p(\boldsymbol{\theta}) p(\mathcal{Z} | \boldsymbol{\theta}) d\boldsymbol{\theta}$ is intractable. The solution is VI (*Variational Inference*)—search a distribution family \mathcal{Q} for the closest distribution to the posterior.⁵ We do this by minimizing the KL divergence between the two,

$$q^* \in \operatorname{argmin}_{q \in \mathcal{Q}} D_{\text{KL}}(q \| p).$$

In our case, we will search the space of isotropic Gaussians. *I.e.*, we search for some mean vector $\boldsymbol{\mu} \in \mathbb{R}^p$ and standard deviation $\sigma > 0$ to form the following distribution,

$$q = \mathcal{N}(\boldsymbol{\mu}, \sigma^2 \mathbf{I}).$$

Thus, we have the following optimization problem,

$$\boldsymbol{\mu}^*, \sigma^* \in \operatorname{argmin}_{\boldsymbol{\mu} \in \mathbb{R}^p, \sigma > 0} D_{\text{KL}}(\mathcal{N}(\boldsymbol{\mu}, \sigma^2 \mathbf{I}) \| p(\boldsymbol{\theta} | \mathcal{Z})).$$

TODO: Justify stochastic gradient descent with Robbins-Monro algorithm.

The Bayesian approach always has three steps—(1) define a prior, (2) define a likelihood, and (3) use Bayes rule to compute the posterior.

⁵ \mathcal{Q} is often called the variational family.

Note that this problem does not have an analytical solution, so we must optimize it with (stochastic) gradient descent. The KL divergence can be rewritten to

$$\begin{aligned} D_{\text{KL}}(q\|p) &\doteq \mathbb{E}_{\theta \sim q} \left[\log \frac{q(\theta)}{p(\theta | \mathcal{Z})} \right] \\ &\propto \mathbb{E}_{\theta \sim q} [\log q(\theta) - \log p(\mathcal{Z} | \theta) - \log p(\theta)]. \end{aligned}$$

Let

$$F(\mu, \sigma, \theta) \doteq \log q(\theta; \mu, \sigma) - \log p(\mathcal{Z} | \theta) - \log p(\theta),$$

then we can derive the gradients to be

$$\begin{aligned} \nabla_{\mu} D_{\text{KL}}(q(\mu, \sigma) \| p(\theta | \mathcal{Z})) &= \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)} [\nabla_{\theta} F(\mu, \sigma, \theta) + \nabla_{\mu} F(\mu, \sigma, \theta)] \\ \nabla_{\sigma} D_{\text{KL}}(q(\mu, \sigma) \| p(\theta | \mathcal{Z})) &= \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)} [\epsilon^{\top} \nabla_{\theta} F(\mu, \sigma, \theta) + \nabla_{\sigma} F(\mu, \sigma, \theta)]. \end{aligned}$$

We can then approximate the gradients by making a single sample Monte Carlo estimate and apply gradient descent to optimize the weights.

7.5 Information-based transductive learning

We are given a domain \mathcal{X} that contains a safe area $\mathcal{S} \subseteq \mathcal{X}$ and area of interest $\mathcal{A} \subseteq \mathcal{X}$. Furthermore, we have an unknown function $f^* : \mathcal{X} \rightarrow \mathbb{R}$ that we wish to explore within \mathcal{A} . For any point in the safe area $x \in \mathcal{S}$, we can query noisy observations,

$$y_x = f^*(x) + \epsilon_x, \quad \mathbb{E}[\epsilon_x] = 0.$$

However, this is prohibitively expensive in some way, so we want to minimize the amount of queries. So, for every query we make, we want to maximize the amount of information that we receive. Formally, given $n - 1$ previous samples $\mathcal{D}_{n-1} = \{(x_i, y_i)\}_{i=1}^{n-1} \subseteq \mathcal{S} \times \mathbb{R}$, we want to find the next point x_n that will give the most information about f in the area of interest \mathcal{A} .

In ITL (*Information-based Transductive Learning*) [Hübotter et al., 2024], the next point is selected as follows,

$$x_n \in \operatorname{argmax}_{x \in \mathcal{S}} I(f_{\mathcal{A}}; y_x | \mathcal{D}_{n-1}).$$

Intuitively, we are looking for a point $x \in \mathcal{S}$ that maximizes the conditional mutual information between y_x and f restricted to \mathcal{A} . When $f \sim \mathcal{GP}(\mu, k)$ with known mean function μ and kernel k , one can show

$$I(f_{\mathcal{A}}; y_x | \mathcal{D}_{n-1}) = \frac{1}{2} \log \left(\frac{\operatorname{Var}[y_x | \mathcal{D}_{n-1}]}{\operatorname{Var}[y_x | f_{\mathcal{A}}, \mathcal{D}_{n-1}]} \right).$$

Proof. TODO: Look at exercises. ■

Next, we will look at special cases of ITL.

We use the reparameterization trick to obtain θ ,

$$\theta = \mu + \sigma \epsilon.$$




Figure 7.1. Space of active learning.

Safe Bayesian optimization. In safe Bayesian optimization, we want to find the maximum of an unknown function f^* , which is done by iteratively choosing points x and observing their realizations $y = f^*(x) + \epsilon_x$. Furthermore, there is an unknown function g^* that defines the safe area,

$$\mathcal{S}^* = \{x \in \mathcal{X} \mid g^*(x) \geq 0\}.$$

We also observe the realizations of this function $z = g^*(x)$. Now, we want to iteratively select points such that we find the maximum of f^* while staying within \mathcal{S}^* .

Assume we are given $n - 1$ points again $\{(x_i, y_i, z_i)\}_{i=1}^{n-1}$. We then fit a GP f on $\{(x_i, y_i)\}_{i=1}^{n-1}$, which induces a lower bound function ℓ_n^f and an upper bound function u_n^f , such that $[\ell_n^f(x), u_n^f(x)]$ is the 95% confidence interval of $\mathbb{E}[f(x)]$ for any $x \in \mathcal{X}$. Similarly, we fit a GP g on $\{(x_i, z_i)\}_{i=1}^{n-1}$, which analogously produces two functions ℓ_n^g and u_n^g . These induce the pessimistic and optimistic estimates of the safe set,

$$\mathcal{S}_n \doteq \{x \mid \ell_n^g(x) \geq 0\}, \quad \hat{\mathcal{S}}_n \doteq \{x \mid u_n^g(x) \geq 0\}.$$

Further, we define the area of interest,

$$\mathcal{A}_n \doteq \left\{x \in \hat{\mathcal{S}}_n \mid u_n^f(x) \geq \max_{x' \in \mathcal{S}_n} \ell_n^f(x')\right\}.$$

Lastly, we can apply ITL using safe set \mathcal{S}_n and area of interest \mathcal{A}_n .

Batch active learning. In batch active learning, we do not select a single point at a time, but we select a batch of points. Formally, we are given an unknown function $f : \mathcal{X} \rightarrow \mathcal{Y}$ that we wish to explore, a population set $X = \{x_1, \dots, x_m\} \subseteq \mathcal{X}$, and a budget $b \leq m$. The goal is to find a subset $L \subseteq X$ such that $|L| = b$ for which we query the oracle $\{f(x) \mid x \in L\}$.

We define $B_\delta(x)$ as the set of the δ -radius ball of points around x ,

$$B_\delta(x) \doteq \{x' \in \mathcal{X} \mid \|x - x'\| \leq \delta\}.$$

We call $B_\delta(x)$ “pure” if f is constant all over $B_\delta(x)$, i.e., every point in a ball around x is classified as the same class. Further, $C(L, \delta)$ is the union of all balls of the chosen subset of points $L \subseteq X$,

$$C(L, \delta) \doteq \bigcup_{x \in L} B_\delta(x).$$

Further, we define

$$C_r(L, \delta) \doteq \{x \in C(L, \delta) \mid \hat{f}(x) = f(x)\}$$

$$C_w(L, \delta) \doteq \{x \in C(L, \delta) \mid \hat{f}(x) \neq f(x)\},$$

where \hat{f} is a classifier of the space \mathcal{X} . Note $C(L, \delta) = C_r(L, \delta) \cup C_w(L, \delta)$. Lastly, we define $\pi(\delta)$ to be the total probability density of points that have an impure δ -radius ball,

$$\pi(\delta) \doteq \mathbb{P}(\{x \in \mathcal{X} \mid B_\delta(x) \text{ is not pure}\}).$$




Figure 7.2. Safe Bayesian optimization.

r stands for “right”.

w stands for “wrong”.

Note that π increases with δ . Suppose we know $\mathcal{Z} = \{(x, f(x)) \mid x \in L\}$, then denote the fitted 1-nearest neighbor classifier on this dataset as \hat{f} .

We have $C_w(L, \delta) \subseteq \{x \in \mathcal{X} \mid B_\delta(x) \text{ is not pure}\}$. (This can easily be verified using Figures 7.3 and 7.4.) So, we have

$$\mathbb{P}(C_w(L, \delta)) \leq \pi(\delta).$$

Furthermore, looking at Figure 7.3, we can easily see the following

$$\{x \in \mathcal{X} \mid \hat{f}(x) \neq f(x)\} \subseteq C_r^c \cup C_w \subseteq C^c \cup \{x \mid B_\delta(x) \text{ is not pure}\},$$

where A^c denotes the complement of a set A . As a result,

$$\begin{aligned} \mathcal{R}(\hat{f}) &\doteq \mathbb{E}[\mathbb{1}\{\hat{f}(x) \neq f(x)\}] \\ &= \mathbb{P}(\hat{f}(x) \neq f(x)) \\ &\leq 1 - \mathbb{P}(C(L, \delta)) + \pi(\delta). \end{aligned}$$

We want to choose L and δ that minimize $\mathcal{R}(\hat{f})$. This is intractable, so we minimize the upper bound instead. This is done by picking δ first and then choosing L that maximizes $\mathbb{P}(C(L, \delta))$. *I.e.*, we want to solve

$$L^* \in \operatorname{argmax}_{L \subseteq X, |L|=b} \mathbb{P}\left(\bigcup_{x \in L} B_\delta(x)\right).$$

This has two problems—the distribution of x is unknown and this problem is NP hard. We address the first problem by approximating the distribution with the empirical distribution induced by X . As a result, we get the following optimization problem,

$$L^* \in \operatorname{argmax}_{L \subseteq X, |L|=b} \frac{1}{|X|} |\{x' \in X \mid \|x' - x\| \leq \delta, \exists x \in L\}|.$$

We circumvent the second problem by greedily picking points with the most other points in its ball that are not within the ball of previously picked points.




Figure 7.3. Batch active learning. The dashed line is the fitted 1-nearest neighbor classifier on this dataset. Here, $C_w(L, \delta)$ denotes the black area and $C_r(L, \delta)$ denotes the gray area.




Figure 7.4. $\pi(\delta)$ is the probability density of the marked area, which is the space of impure data points under a given δ .

8 *Convex optimization*

TODO

8.1 *Support vector machine*

TODO

9 *Ensembles*

TODO

9.1 *Bagging*

TODO

9.2 *Random forests*

TODO

9.3 *AdaBoost*

TODO

10 Stable diffusion




Figure 10.1. Architecture of stable diffusion.

10.1 Diffusion models

An SDE (*Stochastic Differential Equation*) is a differential equation in which one or more terms is a stochastic process, resulting in a solution that is also stochastic. Typically, the SDE of a diffusion process is of the following form,

$$d\mathbf{X}_t = \boldsymbol{\mu}(\mathbf{X}_t, t)dt + \sigma(\mathbf{X}_t, t)d\mathbf{W}_t,$$

where \mathbf{W}_t is a Wiener process (or Brownian motion). This equation tells us that the change in \mathbf{X}_t is driven by a deterministic factor $\boldsymbol{\mu}(\mathbf{X}_t, t)$ and a stochastic factor $\sigma(\mathbf{X}_t, t)d\mathbf{W}_t$. Note that $\boldsymbol{\mu}(\cdot, \cdot)$ and $\sigma(\cdot, \cdot)$ induce a probability distribution over time, p_t of \mathbf{X}_t .

Anderson [1982] showed that the reverse SDE of the diffusion process can be computed as follows,

$$d\mathbf{X}_t = \left[\boldsymbol{\mu}(\mathbf{X}_t, t) - \sigma^2(\mathbf{X}_t, t) \nabla_{\mathbf{X}} \log p_t(\mathbf{X}_t) \right] dt + \sigma(\mathbf{X}_t, t)d\bar{\mathbf{W}}_t,$$

where $\bar{\mathbf{W}}_t$ is a standard Wiener process when time flows backwards from T to 0, and dt is an infinitesimal negative timestep.

Using the DDPM scheduler, diffusion models have the following forward process,

$$\mathbf{x}_{t+1} = \sqrt{1 - \beta_t} \mathbf{x}_t + \sqrt{\beta_t} \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}).$$

Conditioned on x_t , we can reconstruct x_{t-1} as follows with a predicted noise ϵ_θ ,

$$x_{t-1} = \frac{1}{\sqrt{\bar{\alpha}_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) + \sqrt{1 - \alpha_t} z, \quad z \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), \quad (1)$$

where $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{\tau=1}^t \alpha_\tau$. Thus, we need to learn this function.

It can be shown that a diffusion model with the DDPM scheduler is an approximation of a discretization of the following SDE,

$$dx_t = -\frac{1}{2}\beta_t x_t dt + \sqrt{\beta_t} dw_t.$$

The reverse process is thus given by the following reverse SDE,

$$dx_t = \left[-\frac{1}{2}\beta_t x_t - \beta_t \nabla_{x_t} \log p(x_t) \right] dt + \sqrt{\beta_t} d\tilde{w}_t.$$

In practice, we train a diffusion model by randomly sampling $x_0 \sim p_0, t \sim \text{Unif}([T]), \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ and performing a gradient step on the following loss function,

$$\ell = \left\| \epsilon - \epsilon_\theta \left(\sqrt{1 - \beta_t} x_0 + \sqrt{\beta_t} \epsilon \right) \right\|^2.$$

We can sample by iteratively denoising using Equation (1), starting from $x_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$.

Drift term $\mu(x_t, t) \doteq -\frac{1}{2}\beta_t x_t$; diffusion term $\sigma(t) \doteq \sqrt{\beta_t}$. dw_t is a Gaussian with variance dt . We can show that this approximates the diffusion model by discretizing,

$$\begin{aligned} x_{t+1} - x_t &= -\frac{1}{2}\beta_t x_t + \sqrt{\beta_t} \epsilon, \quad \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \\ x_{t+1} &= \left(1 - \frac{1}{2}\beta_t \right) x_t + \sqrt{\beta_t} \epsilon \\ &\approx \sqrt{1 - \beta_t} x_t + \sqrt{\beta_t} \epsilon. \end{aligned}$$

10.2 U-net

U-nets [Ronneberger et al., 2015] are models used for image-to-image translation tasks. In the case of diffusion models, we have such a task, where we get x_t as input and want to predict ϵ . This framework is used to model ϵ_θ . U-nets work by downsampling the input in stages and then upsampling back to the original space in the same stages. At every step of upsampling, the output of the corresponding downsampling step is concatenated to its input. In this way, we get low-level and high-level information.

10.3 Latent diffusion models

Stable diffusion [Rombach et al., 2022] performs diffusion modeling in the latent space of a pretrained VAE [Kingma, 2013]. During training, we thus first map the input image $x_0 \in \mathbb{R}^d$ to its latent encoding,

$$z_0 = \mathcal{E}(x_0), \quad z_0 \in \mathbb{R}^{d'}, \quad d' \ll d.$$

Then, we use the same loss function as above, where we sample a random timestep and noise,

$$\ell = \left\| \epsilon - \epsilon_\theta \left(\sqrt{1 - \beta_t} z_0 + \sqrt{\beta_t} \epsilon \right) \right\|^2, \quad \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_{d'}).$$

Then, during inference, we sample a noise vector in the latent space $z_T \sim \mathcal{N}(\mathbf{0}, I_{d'})$ and denoise using Equation (1) to get \tilde{z}_0 . Lastly, we decode the latent vector back into pixel space,

$$\tilde{x}_0 = \mathcal{D}(\tilde{z}_0).$$

This process requires a well-behaving latent space, so the regularization term that the VAE framework places on the latent space is very important.

10.4 Text embeddings

In order to perform text-to-image generation, we will need a continuous high-dimensional representation of the input text. For this, we use CLIP [Radford et al., 2021]. CLIP trains image and text transformer models to align text-image pairs. Because of this, they contain more semantic embeddings than other methods of training models to obtain text embeddings. Given a text input sequence of size T_c , the CLIP text model returns a sequence of embeddings,

$$\mathbf{C} \in \mathbb{R}^{T_c \times d_c}.$$

We denote this matrix by \mathbf{C} , because we use it for conditioning.

This sequence is contextualized, because CLIP makes use of self-attention.

10.5 Cross-attention

Now the question becomes how to condition a U-net on the input text sequence $\mathbf{C} \in \mathbb{R}^{T_c \times d_c}$. Stable diffusion [Rombach et al., 2022] does this by making use of cross-attention blocks, which it places at the end of every downsampling stage of the U-net. It first rearranges the output of the downsampling block into timesteps,

$$\mathbf{X} \in \mathbb{R}^{T \times d}.$$

Then, it computes queries from \mathbf{X} , and keys and values from \mathbf{C} ,

$$\begin{aligned} \mathbf{Q} &= \mathbf{X}\mathbf{W}_Q, & \mathbf{W}_Q &\in \mathbb{R}^{d_k \times d} \\ \mathbf{K} &= \mathbf{C}\mathbf{W}_K, & \mathbf{W}_K &\in \mathbb{R}^{d_k \times d_c} \\ \mathbf{V} &= \mathbf{C}\mathbf{W}_V, & \mathbf{W}_V &\in \mathbb{R}^{d_v \times d_c}. \end{aligned}$$

It uses these to perform the attention mechanism with a residual connection,

$$\mathbf{\Xi} = \mathbf{X} + \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_k}}\right)\mathbf{V},$$

where $\mathbf{\Xi}$ denotes the conditioned representation of \mathbf{X} . Note that this architecture is agnostic to the type of the condition. As long as we can embed the conditioning variable, we can condition on it in this way—e.g., we can additionally condition on images [Ye et al., 2023].

References

- Brian DO Anderson. Reverse-time diffusion equation models. *Stochastic Processes and their Applications*, 12(3):313–326, 1982.
- Jonas Hübötter, Bhavya Sukhija, Lenart Treven, Yarden As, and Andreas Krause. Information-based transductive active learning. *arXiv preprint arXiv:2402.15898*, 2024.
- Diederik P Kingma. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *International conference on machine learning*, pages 8748–8763. PMLR, 2021.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 10684–10695, 2022.
- Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *Medical image computing and computer-assisted intervention—MICCAI 2015: 18th international conference, Munich, Germany, October 5-9, 2015, proceedings, part III 18*, pages 234–241. Springer, 2015.
- Pragya Sur and Emmanuel J Candès. A modern maximum-likelihood theory for high-dimensional logistic regression. *Proceedings of the National Academy of Sciences*, 116(29):14516–14525, 2019.
- Victor Veitch, Alexander D’Amour, Steve Yadlowsky, and Jacob Eisenstein. Counterfactual invariance to spurious correlations in text classification. *Advances in neural information processing systems*, 34:16196–16208, 2021.
- Hu Ye, Jun Zhang, Sibor Liu, Xiao Han, and Wei Yang. Ip-adapter: Text compatible image prompt adapter for text-to-image diffusion models. *arXiv preprint arXiv:2308.06721*, 2023.